



**PES**  
**UNIVERSITY**

**ML LAB-14**

**Section:- F**

**NAME:- Mohit Kumar**

**SRN:- PES2UG23CS350**

**5TH SEM**

## **Introduction:-**

This laboratory assignment focused on designing, implementing, and training a Convolutional Neural Network (CNN) using PyTorch to classify hand gesture images into three distinct categories: rock, paper, and scissors. The primary objective was to gain practical experience with deep learning fundamentals, including CNN architecture design, data preprocessing pipelines, model training optimization, and performance evaluation metrics. This hands-on project provided valuable insights into computer vision tasks and demonstrated the effectiveness of convolutional neural networks in extracting spatial features from image data for accurate multi-class classification.

## **Model Architecture:-**

### **Overall Architecture Design:-**

The implemented CNN model follows a classic architecture pattern consisting of two main components: a convolutional feature extraction block followed by a fully-connected classification head. This design leverages the strength of convolutional layers to learn hierarchical spatial features while using dense layers for final decision-making.

### **Convolutional Feature Extraction Block:-**

The convolutional block consists of three sequential convolutional layers, each followed by activation and pooling operations. This progressive architecture enables the network to learn increasingly abstract and complex visual features.

#### **Block 1 (Low-Level Features):-**

- Conv2d Layer: 3 input channels (RGB) to 16 output channels
- Kernel Size: 3x3 with padding 1

- Activation: ReLU
- MaxPool2d:  $2 \times 2$  pooling reducing dimensions from  $128 \times 128$  to  $64 \times 64$
- Purpose: Detects edges, corners, and basic patterns

### Block 2 (Mid-Level Features):-

- Conv2d Layer: 16 input channels to 32 output channels
- Kernel Size:  $3 \times 3$
- Activation: ReLU
- MaxPool2d: Reduces  $64 \times 64$  to  $32 \times 32$
- Purpose: Learns textures, shapes, and mid-level patterns

### Block 3 (High-Level Features):-

- Conv2d Layer: 32 input channels to 64 output channels
- Kernel Size:  $3 \times 3$
- Activation: ReLU
- MaxPool2d: Reduces  $32 \times 32$  to  $16 \times 16$
- Purpose: Learns complex hand gesture structures

### Important Design Choices:-

- Kernel Size  $3 \times 3$  used for effective receptive field growth.
- Progressive channel expansion  $3 \rightarrow 16 \rightarrow 32 \rightarrow 64$  allows deeper feature representation.
- MaxPooling reduces computation and provides translational invariance.

Final Feature Map Size:  $64 \text{ channels} \times 16 \times 16 = 16384 \text{ features}$ .

## **Fully-Connected Classifier Block:-**

- Flatten layer converts 3D feature maps into a single vector.
- Linear layer: 16384 to 256 neurons, followed by ReLU.
- Dropout ( $p=0.3$ ) prevents overfitting.
- Output layer: 256 to 3 neurons for class logits.

## **Total Model Parameters:-**

The model contains approximately 4.2 million trainable parameters.

## **Training and Performance:-**

### **Dataset Configuration:-**

Dataset: Rock-Paper-Scissors (Kaggle)

Total Images: 2188

Training Set: 1750 images

Test Set: 438 images

Classes: rock, paper, scissors (balanced)

### **Preprocessing Steps:-**

- Resize to 128×128
- Convert to tensor
- Normalize using mean 0.5 and std 0.5 per channel

## **Training Hyperparameters:-**

- Optimizer: Adam
- Loss Function: CrossEntropyLoss
- Learning Rate: 0.001
- Batch Size: 32
- Epochs: 10
- Random Seed: 42

## **Training Progress:-**

### **Epoch-wise training loss:**

1. 0.7192

2. 0.2184

3. 0.1049

4. 0.0537

5. 0.0267

6. 0.0128

7. 0.0125

8. 0.0041

9. 0.0031

10. 0.0042

#### Observations:-

- Rapid early convergence
- Smooth, stable training
- No signs of overfitting

#### **Test Performance:-**

- Final Test Accuracy: 97.95 percent
- Correct Predictions: 429 out of 438
- Misclassifications: 9

#### **Conclusion and Analysis:-**

#### **Results Discussion:-**

The CNN achieved a strong accuracy of 97.95 percent on unseen data. The architecture depth was adequate and dropout prevented overfitting. Adam optimizer ensured fast convergence. Preprocessing ensured consistent learning behavior.

#### **Challenges Faced:-**

- Dataset path inconsistencies across systems
- Managing GPU memory with batch sizes
- Monitoring convergence and selecting optimal epochs
- File path issues during single-image prediction

## **Suggestions for Future Improvements:-**

- Data Augmentation: flipping, rotation, scaling, color jitter, elastic transforms
- Using Batch Normalization, Residual Connections, or Attention
- Applying Transfer Learning with models like ResNet18 or MobileNet
- Hyperparameter tuning and scheduler usage
- Model ensembling for higher accuracy
- Longer training with early stopping
- Targeted error analysis of misclassified samples